
Second-Hand Moral Knowledge

Author(s): Karen Jones

Source: *The Journal of Philosophy*, Feb., 1999, Vol. 96, No. 2 (Feb., 1999), pp. 55-78

Published by: Journal of Philosophy, Inc.

Stable URL: <http://www.jstor.com/stable/2564672>

JSTOR is a not-for-profit service that helps scholars, researchers, and students discover, use, and build upon a wide range of content in a trusted digital archive. We use information technology and tools to increase productivity and facilitate new forms of scholarship. For more information about JSTOR, please contact support@jstor.org.

Your use of the JSTOR archive indicates your acceptance of the Terms & Conditions of Use, available at <https://about.jstor.org/terms>



is collaborating with JSTOR to digitize, preserve and extend access to *The Journal of Philosophy*

THE JOURNAL OF PHILOSOPHY

VOLUME XCVI, NO. 2, FEBRUARY 1999

SECOND-HAND MORAL KNOWLEDGE*

Trust enters into the making of a virtuous person in at least two ways. First, unless a child has a sufficiently trusting relationship with at least one adult, it is doubtful that she will be able to become the kind of person who can form ethically responsible relationships with others. Infant trust, as Annette Baier¹ has reminded us, is the foundation on which future trust relationships will be built; and when such trust is irreparably shaken, the adult into whom the child grows may be forever cut off from intimacy. Second, a moral beginner must trust other people's moral judgment while she learns to be good. A child does not learn how to be good on her own. She follows the precepts and examples of others, absorbs the moral values of the community into which she has been born, and relies on others who, with luck, will not lead her astray but will turn out to be trustworthy in offering her guidance and correction. While there is some empirical evidence that children learn about morality better when given reasons in support of what they are being taught, part of what a child is being taught is what *counts* as a good reason. Thus, no moral beginner could be in a position to assess adequately the cogency of the reasons she is given until she has progressed some distance in her education.² Until that time, the beginner must rely on other people's judgment.

* This paper was written while I was a fellow at Cornell University's Society for the Humanities; I would like to thank the Society for its support. The paper was presented at *Different Voices: A Conference in Honor of Annette Baier*, held at the University of Pittsburgh in 1995, and to audiences at Cornell, the University of Cincinnati, and the 1996 Australasian Association of Philosophy Annual Conference—each audience contributed to the final shape of the paper. Special thanks to Jennifer Whiting, Lisa Rivera, and Donald Ainslie for written comments on earlier versions.

¹ "Trust and Anti-trust," *Ethics*, xcvi (1986): 231-60.

² Some accounts of moral education (for example, Aristotle's) will give a greater role to trust than others do.

That trust enters into the making of a morally good person should thus be uncontroversial. I want to defend a much more controversial thesis: needing to trust other people's moral judgment is not just a stage we go through on the way to becoming morally virtuous. If the world of value is complex, and if our access to it is shaped by our experiences, then even among the morally mature there will continue to be a significant role for moral testimony and thus for trust. We should not think of our reliance on other people's moral judgment as like training wheels for the novice cyclist, wheels which can and ought to be dispensed with, once we have the know-how and the confidence to ride off on our own. Often, we cannot do as well on our own as we could do if we accepted the moral testimony of others. Moreover, just as borrowing scientific knowledge can enhance our capacity to discover truths about the nonmoral world, borrowing moral knowledge can enhance our capacity to understand the world of value.³

I am going to start by giving a sketch of what I take to be the best arguments *against* testimony as a source of moral knowledge. Next, I give a detailed real-life example which aims to make us suspicious of the case against testimony. At the same time, the example tries to be fair to the considerations against testimony and to show why we might find morally attractive someone who resists such borrowing and refuses such trusting. Armed with the example, I then return to reexamine the initial case against moral testimony. In the final section, I argue that making room for moral testimony brings issues of trust to center stage. The wise truster of moral testimony begins from a default stance of distrust, yet can come to trust and, through trusting, can come to have moral knowledge that would otherwise be inaccessible to her.

I. THE CASE AGAINST TESTIMONY IN MORAL KNOWLEDGE

It has seemed simply obvious to some, such as Robert Paul Wolff,⁴ that there is no place for testimony about moral matters:

³ Throughout this paper, I am going to speak of "moral knowledge" and sometimes even of "moral facts," but this is simply for convenience of expression. While I do think that only a realist account of morality can adequately accommodate our moral practice, it is no part of my brief to argue for that here. Quasi-realists (for example, Simon Blackburn, in *Spreading the Word* (New York: Oxford, 1984)) and norm expressivists (for example, Allan Gibbard, in *Wise Choices, Apt Feelings* (Cambridge: Harvard, 1990)) can each offer nonrealist reconstructions of this sort of talk. Only on a very simple subjectivist view will such talk be illicit. And I take it—admittedly without argument—that such a view is no longer a contender.

⁴ *In Defense of Anarchism* (New York: Harper, 1970); also cited in Gerald Dworkin, *The Theory and Practice of Autonomy* (New York: Cambridge, 1988), p. 52. John Hardwig shows how Wolff is wrong about contemporary mathematics, given the proliferation of mathematical specialties; see his "The Role of Trust in Knowledge," this JOURNAL, LXXXVIII, 12 (December 1991): 693-708.

He [the responsible man] may learn from others about his moral obligations, but only in the sense that a mathematician learns from other mathematicians—namely by hearing from them arguments whose validity he recognizes even though he did not think of them himself. He does not learn in the sense that one learns from an explorer, by accepting as true his accounts of things one cannot see for oneself (*ibid.*, p. 13).

We need to rely on the word of an explorer returning from strange foreign lands because we cannot have exactly *his* reasons for believing what he says. ‘We saw it with our own eyes’ cannot be our reason for believing what he says. Unless we go there ourselves, we must rely on his telling us what he saw.

With morals, it might seem to be otherwise: we give *reasons* for our moral views and these reasons can be shared and so can be evaluated by each of us.⁵ As a crude first approximation, to be refined later, *testifying* that *p* contrasts with *arguing* that *p* insofar as it is the testifier herself who vouches for the truth of *p*; someone who argues that *p* lets the arguments vouch for themselves.⁶ So long as we can give arguments for our moral positions, it seems that we do not *ourselves* have to vouch for the truth of those positions. On this view, moral testimony can have a significant place only for beginners who cannot manage to evaluate the arguments. For the rest of us, moral testimony can be useful at most only when we do not have enough time to give careful consideration to all the arguments relevant to a moral position. We might, then, because of time constraints, rely on the moral testimony of those whose character we trust and who we know have been able to give careful consideration to the arguments.

Three additional considerations seem to roll back even further the place for testimony about moral matters.

First, moral matters are important and engage our responsibility in a way that nonmoral matters do not. We are responsible for avoiding errors here in a way that we are not responsible for avoiding them in other areas. Further, moral judgments reflect on our character in ways that other judgments do not. The extra importance attached to moral judgments derives both from their close connection

⁵ C. A. J. Coady mentions this as an important consideration against testimony regarding moral matters; see his *Testimony: A Philosophical Study* (New York: Oxford, 1992), p. 71.

⁶ This account of testifying is not quite the same as that offered by Coady. He builds speaker competence into his definition of testifying (*op. cit.*, p. 43). Elizabeth Fricker rightly takes him to task for this, since it shifts the interesting epistemological questions to the question of whether someone was in fact testifying that *p*—“Telling and Trusting: Reductionism and Anti-reductionism in the Epistemology of Testimony,” *Mind*, civ (1995): 393-41.

with action and from the importance of the concerns that morality addresses.

Second, moral knowledge is supposed to be *practical* knowledge. It is not enough simply to know, for example, that sexism is wrong. One must be able to put that knowledge into practice.⁷ But, insofar as one must rely on other people's judgment as to when 'sexism' applies, one's knowledge of the wrongness of sexism will be practically deficient. Similarly, unless one knows not only *that* sexism is wrong but also *why* it is wrong, one will not be able to generalize from the wrongness of sexism to the wrongness of other sorts of discrimination. Because of this, one's knowledge that sexism is wrong will not reverberate throughout one's actions and will not extend itself to other practical spheres.

Third, borrowing moral knowledge appears to be in tension with the ideal of an autonomous moral agent. On some accounts of autonomy, the tension will be straightforward. If we think of autonomy as requiring independence from others, then those who borrow moral knowledge must, by that fact alone, fail to be autonomous.⁸ On this picture, being autonomous presupposes what we might call "intellectual individualism."⁹ Such accounts of autonomy have been roundly criticized, for it is hard to see why autonomy, conceived as independence, should be something valuable.

A more plausible account of autonomy views autonomy as involving, among other things, the capacity to reflect critically on the principles that one adopts. Through such reflection, one comes to make those principles one's own.¹⁰ But it is easy to see why this account of autonomy, while not as obviously incompatible with borrowing moral knowledge as autonomy-as-independence views, might nonetheless be in tension with it. If we have borrowed our moral knowledge from others and cannot offer reasons in support of it, then how can we reflect critically on that knowledge? Moreover, if we do not understand the rationale behind those borrowed principles, how can we use our borrowed principles to reflect critically on other principles we accept? How do we know the ways in which those

⁷ This point is made by G. E. M. Anscombe, in "Authority in Morals," reprinted in *The Collected Philosophical Papers of G. E. M. Anscombe* (Minneapolis: Minnesota UP, 1981), pp. 43-50, here pp. 47-48.

⁸ For this view, see Wolff; and Thomas Scanlon, "A Theory of Freedom of Expression," *Philosophy and Public Affairs*, 1 (1972): 204-26. For cogent criticisms of it, see Dworkin, pp. 22-25.

⁹ This term is from Sigurdur Kristinsson's Ph.D. dissertation, *The Nature and Value of Individual Autonomy* (Cornell University, 1996), p. 129.

¹⁰ Dworkin gestures toward such a requirement (*op. cit.*, p. 20).

principles connect with others, or the ways in which the considerations that support them align or fail to align with the considerations that we take to support our other principles? And if we cannot bring our borrowed knowledge into play in rationally assessing our other moral commitments, then those commitments will not be given the critical scrutiny they deserve.

We began by thinking that since moral views are held for reasons, moral argument should be our primary source for moral knowledge. Testimony can have a place, though a small one, because we do not always have time to examine adequately the reasons in favor of a moral view. Now it looks as though it is incumbent on us, as conscientious agents, to limit our dependency on the moral testimony of others as much as we can. If, in the end, we sometimes still have to rely on testimony, that is a regrettable consequence of our finiteness.

II. AN EXAMPLE THAT REOPENS THE CASE AGAINST MORAL TESTIMONY
The case against testimony about moral matters appears to be strong. It is easy to construct examples that show, however, that the case against testimony must be less than fully compelling. Here is one such example.

Peter had lived in a fourteen-person cooperative house for two years and was deeply committed to the ideal of cooperative living. His first year there, he had felt uneasy about the cooperative's membership decisions, but the second year he found them intolerable. Membership decisions were made on the basis of quite lengthy interviews, and decisions to accept someone had to be unanimous. Any member could exercise a veto over a potential member for any reason, though in practice such vetoes were rarely necessary, as support for a potential member who might have incurred a veto generally evaporated. Peter came to believe that the threat of veto was making membership decisions arbitrary and unjust. White women and women of color had advocated rejecting three white men on the basis of their perceived sexism, and, in one case, racism. Peter had a settled and serious commitment to the elimination of racism and sexism, but he was not very good at picking out instances of sexism and racism. Thus, his commitment often remained theoretical and he saw fewer occasions as calling for antisexist or antiracist response than he might have seen. Such blindness can sometimes indicate insincerity, but in Peter's case it did not. He genuinely did want to understand the perspective of the women and he wanted to be able to share it, *if* indeed it turned out to be correct. He could pick out egregious instances of sexism and racism, and could sometimes see that "sexist" or "racist" applied to more subtle instances when the

reason for their application was explained to him, but he seemed bad at working out how to go on to apply the words to nonegregious new cases. Problems arose for Peter when he could not grasp the reasons why the women were calling someone sexist, when he either could not see, or could not see as evidence, the considerations that the women thought supported viewing the would-be members as sexist. Discussion tended to trail off:

"No, it wasn't anything he said, it was the way he looked at me when I asked him a question."

"How do you mean?"

"Well, it wasn't that he didn't look at me and he didn't stare at my breasts, it was just a way of looking—a way of not quite being present as he answered me."

"What way? I do not understand what way you mean."

"Well, it was just a feeling I got. A feeling that he didn't think me important."

"But you can't reject someone for a feeling."

"I'm not, I'm rejecting him because he's sexist."

Peter decided to leave the co-op. His reason for leaving was that, as a member, he felt he had to be able to endorse the decisions that the group was making—they had to be decisions that were "his" in the sense that he could identify with them, even if they were the result of compromise and negotiation.

The women in the house wanted to say, and did say: "Look, just trust us, we know about this." But to Peter, such trusting was an abdication of moral responsibility. Even worse, from Peter's point of view, the decisions about which men to accept were arbitrary and discriminated against the men rejected. Nor could such discrimination be justified on the basis of some goal it might achieve; it was not, for instance, a matter of increasing the presence of women and people of color in the co-op. It was, from his point of view, prejudice, pure and simple.

I think that in this case Peter should have been willing to accept the women's testimony that these men were sexist. But for him to have done so would have been to accept and act upon a moral judgment on someone else's say-so. It would have been to borrow moral knowledge.

At this point, it might be objected that the case is misdescribed. Peter was not being asked to rely on other people's moral judgment at all. Instead, he was being asked simply to accept that a purely descriptive term, 'sexist', applied. If we are inclined to think it not *outrageous* that he should have deferred to the women of the house, that

is because it is reasonable to suppose that they might be better able than he to apply the descriptive terms ‘sexist’ and ‘sexism’. But, the objection continues, this is no different from our deferring to biologists over when species terms apply, to astronomers over astronomy terms, and so on. It is true that ‘sexism’ also carries with it connotations of disapproval, but we can divide the meaning of the term into an expressive or condemnatory component and a purely descriptive component. The most popular candidate for the descriptive component would be ‘differentiating between the sexes on the basis of their sex and nothing else’.

This objection is mistaken, however. ‘Sexism’ refers not to just *any* sort of differentiation between the sexes, but to morally unjustifiable differentiation. (It is a failure to grasp this that leads some to say that *any* women-only event has to be sexist since it treats men and women differently simply on the basis of sex.) Thus, to know how to apply the term ‘sexism’ is to have some moral knowledge.¹¹

The objection does point out an important feature of the case, however, and one that shows it to be the simplest, and most readily defended, sort of case. Peter already *had* a commitment to ending sexism. He was not borrowing from others the very principle “sexism is wrong.” That much, let us suppose, he had learned *for himself* (though probably not *by himself*). If asked, he would have been able to give reasons in support of it—he might, for example, have said that sexism shows a failure to accord all persons equal respect. Peter did not need to borrow his knowledge of the relevant principle. What he needed to accept on other people’s say-so was that this principle applied. We can thus distinguish between taking it on trust that a principle you independently acknowledge applies in a particular

¹¹ Sometimes there is a clear answer to the question of whether something is moral knowledge or the sort of nonmoral knowledge that feeds into moral knowledge. (For example, average daily caloric intake of the poor may be the sort of factor on which the justice of a society supervenes, but it is clearly nonmoral.) But I do not think this distinction is always easy to draw. For example, should we say that the fact that someone will be upset by a course of action is a purely nonevaluative fact, though one which may contribute to determining whether a course of action is right or wrong? Or should we say it is itself an evaluative consideration? There is reason for thinking it is evaluative. Being upset contrasts with what is colloquially called “having a hissy fit,” since the later implies that the distress is *prima facie* unjustified. ‘Upset’ may thus have connotations of *prima facie* justified (or at any rate not *prima facie* unjustified) distress. It is worth noting that, if you want to maintain a sharp separation between moral and nonmoral epistemology and allow testimony to the latter but not the former, then you had better be able to draw clear lines between the moral and the nonmoral. A view like my own, that postulates no such sharp dichotomy at the level of epistemology, need not worry about being able to draw this distinction in every case.

case, and taking it on trust that a principle is correct. In between these two, we can also identify taking it on trust that the principles you recognize on independent grounds are to be ranked in a given way.¹² (Such rankings are often at issue in setting political agendas.) We shall need to consider whether there is reason to be more or less suspicious of borrowing moral knowledge depending on what kind of knowledge it is that you are borrowing.

While I think that Peter was mistaken to refuse to accept the women's testimony regarding whether someone was sexist, his refusal nonetheless displayed admirable traits of character and admirable strength of character. Peter's action was not motivated by a desire to control the outcome of collective decisions. He made the decision to leave quietly and it was never intended as a threat to force the group to decide in the way he thought correct. The esteem in which Peter and the women held each other made it possible to sustain friendships despite the disagreement and despite Peter's leaving. It was quite probable, too, that others shared Peter's misgivings, but lacked his conscientiousness about moral matters.

Peter's stance shows us the *prima facie* plausibility of the three considerations against accepting moral testimony: he shows a keen and admirable awareness of the importance of moral issues. Moreover, his desire to have, *and to understand*, morally adequate reasons for acting makes us fairly sure, that, should Peter fall into morally bad circumstances, he will not be following the leader into error. This sort of independence—the independence of those who want to keep their moral beliefs in step with what they take to be cogent reasons—gives us a sense of the notion of autonomy which shows why autonomy is thought to be something valuable. Finally, had Peter accepted the women's testimony, he would still not have been any the wiser about how to go on to apply 'sexism' to new, nonegregious, cases.

Overall, though, Peter's case does more to undermine the arguments against testimony than to support them. In particular, it in-

¹² This tripartite distinction does not map neatly onto Gibbard's distinction between fundamental and contextual authority (*op. cit.*, pp. 174-81). Contextual authority is authority granted to another on the presupposition of shared norms. Gibbard's examples, however, suggest that he may have in mind cases in which the borrower is ignorant about some nonmoral facts of the situation and so lets someone else's moral reasoning stand "proxy" (*op. cit.*, p. 174) for her own. A moral agent claims fundamental authority when she claims the superior trustworthiness of her own norms. But it seems that one can accord another *some* fundamental authority without accepting their norms on the basis of her say-so, as when one takes the fact that someone believes a norm as a reason for accepting it, but as a reason that is, in the circumstances, outweighed (*op. cit.*, p. 180).

vites us to reexamine the initial assumption that led us to believe that testimony must, for the morally mature, always be an inferior way of learning about morality relative to learning about it through argument. We are not supposed to need testimony, except when time is short, because someone who accepts a moral position can always make the grounds of her beliefs accessible to us, and we can come to share her reasons or not, depending on whether we find them persuasive.

Reflection on Peter's case suggests two compatible responses to this argument. We might say that at least sometimes, the capacities involved in acquiring moral knowledge are indeed similar to perceptual capacities. There is a moral analogue of blindness, which is not to be assimilated to being too foolish to follow a moral argument. Those who lack the relevant perceptual skills in a given moral domain must rely on those who have sharper moral vision.

Alternatively, we might point out that the contrast between testifying and arguing is not hard and fast. Expert witnesses both testify and argue. They offer arguments when presenting the considerations that lead them to adopt their views, but they testify, too, since the nonexpert to whom the testimony is directed is not in a position to assess fully the cogency of the reasons the experts offer in favor of their views. The nonexpert may not know exactly why a certain response on a given test indicates with such-and-such probability that the blood sample comes from the defendant, but will, having established the expert's credibility, take her word that it does. Offering arguments for a position thus need not entirely eliminate the reliance on the word of the testifier which is characteristic of offering testimony. Expert testimony shows the context sensitivity of the notion of testimony—directed at fellow experts, the arguments that for the lay person are also testimony would be just arguments.

I think that both responses have a place and are not as problematic as they might first appear. In defense of the claim that moral knowledge sometimes requires skills analogous to perception, we need only think of those occasions in which seeing that some moral value is at stake requires an emotional sensitivity to the particulars of a situation. For example, seeing that a situation calls for kindness may require an empathetic response involving imagination of what it must be like to be in someone else's shoes. But we might not all be equally capable of such empathetic leaps—some of us may need to rely on those who are.

Sometimes, those from whom we accept moral testimony stand to us, regarding certain issues, as moral experts. But the very idea of moral experts might seem deeply problematic. It seems profoundly

inegalitarian to claim that some are moral experts, since it suggests that others are not.¹³ While the higher branches of mathematics might be beyond some of us regardless of our best attempts, and available to others only if they make certain attempts, it is often thought that each of us is equally capable of being a good moral agent. And there can be no experts where capacity is equal.¹⁴ Moreover, an induction on the past history of those who have claimed to be moral experts and with what results might make us skeptical of the very idea of moral experts. Many who have claimed that role have used it to advance self-serving and morally dubious views.

These concerns about moral experts need not be taken as concerns about the very idea of moral expertise. They are for the most part generated by worries about establishing a socially sanctioned role of moral expert. Having a socially sanctioned role of moral expert is going to raise issues of credentialing—Who are the moral experts and who gets to decide?—and issues of what authority those in the socially sanctioned role of moral expert are to have over policy formation, and over the rest of us who may disagree with their judgments. These issues, however, do not arise when considering the question of what sorts of circumstances, if any, are ones in which an agent should decide to rely on someone else's moral judgment.

We can further reduce the sting of acknowledging that there can be expert testimony on moral matters by distinguishing between *experts* and *expertise*,¹⁵ and between *global* and *local* claims to expertise. While we think that there can be few experts on any subject—if your knowledge is not rare and hard to come by, you would not merit the title “expert”—expertise can be widely shared. Normal adults have considerable expertise at a wide variety of things. To borrow a list

¹³ See Arthur Caplan, “Moral Experts and Moral Expertise: Does Either Exist?” reprinted in his *If I Were a Rich Man Could I Buy a Pancreas?* (Bloomington: Indiana UP, 1992), pp. 18-39.

¹⁴ This thought lies behind the strong suspicion that if, for example, Lawrence Kohlberg’s scale of moral development consistently scores women, the working class, and members of other cultures at a lower level of moral development than educated middle-class white males, this must indicate some bias in the scale. The suspicion of bias is not removed by the fact that when education and type of employment are controlled for, men and women score the same. In contrast, we would not infer bias in a scale that aimed to measure mathematical skill if it led to similar results. We would simply say that those with certain educational backgrounds and occupations were, most likely, better mathematicians than those who lacked these backgrounds. While it might be plausible that having a certain background exposes you to certain *types* of moral problems, however, everyone, we think, has exposure to moral problems of some type or other, and thus everyone will have the sort of experience necessary to develop their potential as moral agents, though this development may take different forms.

¹⁵ For this distinction, see Caplan.



from Arthur Caplan, "how to drive a car, mail a letter, exhibit good manners, use a post office, play games" (*op. cit.*, pp. 30-31), and, we might add, considerable expertise in ethical matters. Thus, in saying that someone can sometimes stand to another much as expert witness to lay person, we are not committed to saying that the person who currently does not have the expertise could not acquire it, though in some cases the kinds of character change and changes in experience necessary in order for them to acquire it may be sufficiently radical that, for all practical purposes, they are unable to do so. More important, we are not committed to saying that one person stands to another as expert to lay person in *all* areas of value. If that were indeed the case, it would be tempting to say that the one forced to rely on the other would lack moral maturity, so we would be talking not about a case in which the morally mature agent needed to borrow moral knowledge but about a case in which the agent was still a beginner.

Although lowering the cost of the analogy between expert testimony and testimony regarding morality, none of what has been said so far gives us any reason to suppose that moral expertise might be unevenly distributed with respect to certain domains.¹⁶ But evidence for this claim is not hard to find. Examples like that of Peter provide us with the sort of evidence we need. Nor does positing such differences in likely familiarity and knowledge commit us to any epistemologically radical claims. Baier (*op. cit.*, p. 247) and others argue that one's moral experience shapes the kinds of questions one will find salient: had moral theorists begun from the experience of women, they would have paid more attention to trust, since trust is of vital importance in relations among unequals.

Political engagement in liberation movements can also be expected to lead to greater insight and understanding about the values at stake than is likely to be available to someone who lacked that sort of experience. One way to find out about the value of respect is to ask those who have always been respected; a better way is to ask those who have struggled to win respect and tried to live without it. Saying that such political experience can contribute to moral knowledge commits us to nothing stronger than the view that one's knowledge about moral matters, like one's knowledge about other matters, can deepen with experience. In particular, the suggestion that, for example, women actively engaged in feminist movement, might have acquired some expertise about the

¹⁶ In fact, this would have to be concerning certain areas only, or we would run afoul of the idea that, globally, there will not be much difference in moral skill between people in distinct groups.

kind of disvalue sexism is, and the often subtle forms it can take, does not require us to defend a notion of epistemic privilege of the kind defended by standpoint theorists.¹⁷ It is enough if, through social location or choice, individuals come to have richer experiences with certain types of moral problems. We need not follow standpoint theorists and suppose in addition that the marginalized are always, or even generally, in the best position to understand the truth about social relations.¹⁸

If we allow that capacities akin to perceptual capacities are sometimes involved in generating moral knowledge—especially the knowledge of when values are at stake in particular situations—and if we admit that there can be local differences in moral expertise, differences explained by different experiences, then we have come some way toward showing that there can be room for moral testimony. Sometimes, we do need to hear, and to believe, reports of explorers returning from realms of value that are less familiar to us than to them. Nonetheless, one might continue to be uneasy about moral testimony: What of the objections that borrowing moral knowledge underestimates the importance of moral matters, overlooks their practicality, and conflicts with autonomy? The first of these objections is best dealt with now, while the second and third need to wait until we have a clearer understanding of when an agent is justified in relying on moral testimony.

We need not suppose that we have to meet the first objection on exactly its own ground. There is a good case to be made for just denying that there is any difference between moral and other knowledge with respect to the way it engages our responsibility: we are every bit as accountable for believing the racist science of *The Bell Curve*¹⁹ as we are for having false moral beliefs. But the strongest case against this position can be made by affirming its premise. If it is so

¹⁷ For this use of ‘feminist movement’ rather than ‘the feminist movement’, see bell hooks, *Feminist Theory from Margin to Center* (Boston: South End, 1984).

¹⁸ For a discussion of standpoint theory, see Nancy Hartsock, “The Feminist Standpoint: Developing the Ground for a Specifically Feminist Historical Materialism,” in Sandra Harding and Merrill B. Hintikka, eds., *Discovering Reality: Feminist Perspectives on Epistemology, Metaphysics, Methodology, and Philosophy of Science* (Dordrecht: Reidel, 1983), pp. 283-310. It would be a mistake to think that political engagement must count against objectivity. Even if such engagement means that one will approach questions with presuppositions, provided that one’s presuppositions are true (or approximately true), one can count as being objective. Note that unless this were the case, the theory dependence of scientific method would count against the objectivity of science. For this notion of objectivity, see Richard Boyd, “How to Be Moral Realist,” in G. Sayre-McCord, ed., *Essays on Moral Realism* (Ithaca: Cornell, 1988), pp. 181-228; and Louise Antony, “Quine as Feminist: The Radical Import of Naturalized Epistemology,” in Antony and Charlotte Witt, eds., *A Mind of One’s Own* (Boulder: Westview, 1993), pp. 185-225.

¹⁹ Richard J. Herrnstein and Charles Murray, *The Bell Curve: Intelligence and Class Structure in American Life* (New York: Free Press, 1994).

very bad to make a moral mistake, then it would take astonishing arrogance to suppose that this supports a do-it-yourself approach, especially once we think about how expertise can vary with experience. Of course, the worry behind this objection might be that people will too readily rest content with borrowing the moral beliefs of others without any grounds for supposing them likely to be correct. And that, the objection goes, would be a failure to take moral matters as seriously as they should be taken. This assumes, however, that the borrower of moral knowledge has few responsibilities to meet before she can count as borrowing wisely. In what follows, I shall argue that the standards for wise borrowing are stringent indeed and so offer no comfort to the morally indolent.

III. WISE TRUSTING

I shall argue that the problem of how to be a wise receiver of moral testimony is a special case of the problem of how to be a wise receiver of testimony more generally, and this in turn is a special case of the problem of how to be wise in one's trust. Because of this, theories of trust, and in particular accounts of trust's justification conditions, can offer insight into some of the epistemological issues raised by testimony. But first it is necessary to show that relying on testimony does, in the typical case, require trust in the testifier.

It might seem simply obvious that testimony involves trust: where there is testimony there is trust, or so says common sense. The language in which we describe cases of testimony supports this claim, for we speak of "taking it on trust," and of "trusting testimony." While I shall argue that common sense is basically right here, it does take some additional argument to vindicate it—unless, that is, we have a minimalist account of trust.

So as not to mislead the reader into thinking the connection between trust and testimony follows just as a matter of course, I shall temporarily abstain from using phrases like 'trusting testimony' or 'taking it on trust', and use instead the phrase 'accepting testimony'. Someone accepts the testimony of another if she forms the belief that p on the basis of the other's saying that p is the case. We can use testimony that we do not accept: for example, the movie critic whose rave tells us that we would hate the movie, or the suspect whose answer signals that she has something to hide. The thesis that needs to be defended can be put like this: in the typical case, accepting testimony requires trust in the testifier and, as often as not, in others as well.

How plausible one finds the claim that there is a connection between trust and accepting testimony depends, not surprisingly, on

the account of trust one accepts. If you think that trust is just a matter of relying on others, then it is obvious that accepting testimony has got something to do with trust. Call accounts of this kind *risk-assessment accounts*. Diego Gambetta²⁰ offers us an example of a risk-assessment account:

...trust (or, symmetrically, distrust) is a certain level of subjective probability with which an agent assesses that another agent or group of agents will perform a particular action, both *before* he can monitor such action (or independently of his capacity ever to be able to monitor it) *and* in a context in which it affects *his own action* (*ibid.*, p. 217).

On a risk-assessment account, it follows straightforwardly that accepting testimony requires trust: I would not accept your testimony unless I assigned a sufficiently high degree of probability to your telling me the truth in these circumstances.

Risk-assessment accounts do not distinguish among the reasons one might assign a sufficiently high degree of probability to an agent's performing an action. One might think an agent will perform that action out of fear or stupidity, because it coincides with her own self-interest, or because she has, and wishes to display, goodwill toward those who are counting on her.

Will-based accounts, in contrast, find trust only where there is reliance on the goodwill of another. We are not to take 'goodwill' as synonymous with 'personal friendly feeling', or will-based accounts would be obviously too narrow. Sometimes those we trust have the relevant goodwill just in virtue of being morally decent, or honest, or caring about fulfilling their duty. Since will-based accounts are more restrictive than risk-assessment accounts, if I can show that even on such accounts trust is needed, then I shall have shown that *whatever* account of trust is ultimately endorsed, accepting testimony requires trust. Consider the following will-based account of trust according to which trust is:

- (a) an attitude of optimism that the goodwill and competence of another will extend to cover the domain of your interaction, together with
- (b) the confident expectation that the one trusted will be directly and favorably moved by the thought that you are counting on her.²¹

There are clearly some nonstandard cases in which I can accept testimony without having an attitude of optimism that the goodwill of

²⁰ "Can We Trust Trust?" in Gambetta, ed., *Trust: Making and Breaking Cooperative Relations* (New York: Blackwell, 1988), pp. 213-38.

²¹ For a defense of this account, see my "Trust as an Affective Attitude," *Ethics*, cvii (1996): 4-25. Baier defends a will-based account (*op. cit.*).

another will extend to cover the domain of our interaction. I may, for example, have evidence that you would not dare lie because I have a threat advantage, and know that you are someone who is susceptible to threats. In a different scenario, I might know that it will somehow suit your devious purposes to tell me the truth, here and now (though in general you would not). Usually, however, we are in neither position: we have no threat advantage sufficient to guarantee performance without also relying on goodwill, and we do not have the kind of information that would let us know when those who are generally untrustworthy are in fact telling the truth. The false witness, like the lying promise giver, is threatened with the possibility of never again having her testimony believed.²² But this threat is not sufficient to guarantee performance in such a way as to eliminate the need for reliance on goodwill, and to be effective at all it requires both that we be able to tell when we have been deceived and that we will again encounter the false witness.²³ These conditions are frequently not met.

What about the need for optimism about the competence of another? In accepting testimony, we typically lack the kind of proof about the competence of another that would securely ground our optimism in her competence. Indeed, certain kinds of proof are ruled out. If testimony is to be at all useful, I cannot myself go and check up on what you are saying. Sometimes, my not doing so is merely a matter of convenience: I accept testimony precisely because it would be too bothersome to learn for myself what I can learn from you. But where testimony has the potential to have rich epistemic yield, I am not in a position to do the kind of checking up on your judgment which would be required. I may be ignorant about, for example, some mathematical theorem, and lack the expertise to find out the truth of the matter on my own. Even if I attempted to check up on your word by trying to prove it on my own, I, novice that I am, shall often have fewer grounds for confidence in my own judgment than I have for confidence in yours.²⁴

At this point, it might be objected that I have failed to distinguish two kinds of checking up: I may not be able to check up on the *content* of what you say, but I may be able to check up on *you*.²⁵ We often do call on other people's assessment of our informant's character

²² For this penalty in the case of the false promisor, see David Hume, *A Treatise of Human Nature*, P. H. Nidditch, ed. (New York: Oxford, 1978, 2d ed.), p. 522.

²³ For this response to the objection that trust is not needed in the scientific context because game-theoretic considerations alone will be sufficient to guarantee honesty, see Hardwig, p. 704.

²⁴ Hardwig makes this point (*op. cit.*, pp. 698-99).

²⁵ These are distinguished by Dworkin (*op. cit.*, p. 56).

and expertise before accepting their testimony. But this move does not eliminate trust; rather, it grounds my trust in you in my trust in others. It may disperse trust, but it does not make it redundant.

To show that it is true on the account of trust in question that accepting testimony typically requires trust in the testifier, we also need to show that such acceptance will require the confident expectation that the one trusted will be directly and favorably moved by the thought that one who trusts is counting on them. It is easy to show that this is also required: conversational exchanges do not come labeled, conveniently announcing themselves as the kinds of exchanges that they are. Before I can accept what you say in response to my request or need for information, I must first assume that what you are doing is *in fact* responding to my request or need—that you are not, for example, teasing me by saying something silly since you think a silly question deserves a silly answer.

The common-sense assumption that accepting testimony involves trust is borne out, if not for every possible case, then for the standard cases and these form the vast majority of cases of accepting testimony.

If this is right, then the question of what it is to be a wise and responsible recipient of moral testimony just will *be* the question of what it takes to trust wisely in these domains. And that is a nice result, because we know something about the justification conditions of trust.

In *Testimony*, C. A. J. Coady tries to argue that our default stance should be one of trust in testimony: we require some evidence to move away from this stance and none to adopt it. Viewed from the perspective of accounts of when trust is justified, this claim invites suspicion. It is simply impossible to say what our default stance with respect to trust in general should be. For this reason, we should be surprised indeed to find that one could defend a default stance of trust with respect to testimony, given that testimony can cover a wide range of areas, be given by a wide range of testifiers, and occur in different social climates and different institutional settings.²⁶

²⁶ Fricker notes the context sensitivity of appropriate trust in testimony (*op. cit.*). She accuses Coady of an inconsistency in saying that our default stance should be one of trust and at the same time allowing that we accept testimony only when signs of untrustworthiness are absent (*op. cit.*, p. 405). But this is still to have a default stance of trust. So long as I do not need positive evidence of trustworthiness before I trust, then my default stance is one of trust. In what follows, I argue that sometimes our default stance can be distrust; thus my position may be more skeptical than Fricker's.

The reason there are no useful generalizations to be made concerning whether trust, distrust, or neutrality is the appropriate default stance is that the appropriate default stance is sensitive to climate, to domain, and, linked to domain, to consequences. Further, for the individual truster, the appropriate default stance is linked to the reasonable assessment of the tendencies of her own trusting and distrusting.²⁷

In climates in which there is strong motive to be untrustworthy, it would require more evidence to believe that a person was trustworthy than it would take to believe that a person is trustworthy in climates where there is little incentive to untrustworthiness. Thus in unfavorable climates, our default stance should be distrust, and it should take considerable evidence to move from that position to trusting; in more favorable climates, however, a default stance of trust can sometimes be warranted.

Domain also affects which default stance is appropriate: if I am to trust you with a secret of mine, then I shall want to have quite a bit of evidence about your character. If, though, I am to trust you not to attack me in the street, I may need no particular evidence about your character at all. This might seem to be a counterintuitive result, since surely it is worse to be attacked in the street than it is to be embarrassed by a confidence indiscreetly betrayed. Domain is generally more important than the consequences of misplaced trust, however, because domain signals likelihood of performance. We are all aware of the lively attractions of gossip and the strength of character required to resist them. In contrast, it is not hard to refrain from injuring a stranger on the street; that just takes basic decency, a trait that we can assume is widely shared, unless the climate is sufficiently bad. Once we hold domain fixed, consequences become of the first importance: of course, I am going to need more evidence of your trustworthiness before I tell you a secret that, if spread abroad, would be very damaging to me, than before I tell you a secret whose disclosure would be merely embarrassing.²⁸

²⁷ These conditions for justified trusting are further elaborated and defended in my "Trust as an Affective Attitude," especially at pp. 20-25.

²⁸ Hardwig wonders about the extent to which the truster of scientific testimony needs to have good evidence for her belief that the inquirer is trustworthy (*op. cit.*, p. 699). It seems to me that his account would be more plausible if he were to take differences in domain into account. Scientific inquiry has an historically poor track record concerning politically sensitive topics, such as race and gender. Whatever we are to say about the kinds of evidence needed before accepting the research of other scientists in general, it seems we need extra caution when it comes to research that can validate social relations of dominance and subordination.

While climate, domain, and consequences are variables determining which default stance is justified that extend across agents, the fourth variable determining the appropriate default stance is agent specific. Some agents have reason to be distrustful of their tendencies toward trust in certain domains. Others have reason to be distrustful of their tendencies toward distrust. Here we might think of, for example, racist tendencies to distrust young African-American men. Such agents should be suspicious of their own default stances, and that, in turn, should lead them to be more willing to abandon them than those who can trust their trust and distrust.

If climate, domain, consequences, and metastances of trust or distrust are the important variables determining how much evidence is warranted before trusting, we can see why there can be no answer as to the appropriate stance to take toward testimony in general. Most people can be trusted to give an honest and informed answer to certain questions, including the mundane sorts of questions we often rely on strangers to answer for us, such as what time it is, how we get to the bus stop, and so on.²⁹ But when it comes to matters where there are reasons for not being straightforward—for example, social-science surveys of sexual activities—or to matters where competence cannot be assumed, we should not adopt a default stance of trust in testimony.

This general discussion has implications for what we should say about how much and what sort of evidence we should demand before trusting someone's moral testimony. Moral knowledge regarding the sorts of matters for which we might need moral testimony cannot be easy to come by. If it were easy to come by, then we would expect it to be in the commonly shared stock of moral expertise. In addition, our moral thinking is vulnerable to distortion due to, for example, wish fulfillment and self-interest. Having such knowledge requires a good character, as well as the right sorts of experiences. Since the best way to convince others to go along with your interests is to convince them that morality requires them to do so, we can expect tendencies toward untrustworthiness—perhaps deliberate, perhaps as the result of self-deception.

If this is right, then the appropriate default stance toward testifiers about morality is one of distrust. A default stance of distrust gives a presumption against accepting moral testimony, where this pre-

²⁹ But even this list is climate sensitive. I am told that in Egypt it is thought disrespectful not even to attempt to offer directions to those who ask for them. Rather than be disrespectful, an informant will send you in completely the wrong direction.

sumption can take the form of either disbelieving such testimony, or withholding judgment. What distinguishes a default stance of distrust from a stance of neutrality is that a substantial weight of evidence is required to rebut the presumption; with neutrality, there is no presumption that stands in need of rebutting.

In discussing Peter's case, I conceded that it was the simplest sort of case, since it involved relying on someone else's judgment that a principle he already accepted applied. If your source is wrong about the relative importance of a good, or wrong about which principles or ends to adopt, *and you take her word for it*, your error is going to surface not just in one place but in many; if, however, you rely on someone's judgment that a principle you accept on independent grounds is at stake in a given situation and she is wrong, at least the error you inherit is contained. The intuition that these kinds of cases are importantly different is shown to be just a special instance of the way in which consequences interact with domain to affect the justification of trust.

But acknowledging that the appropriate default stance toward testifiers on moral matters is distrust does not mean that we can never be licensed in moving from distrust to a position of trust. Before we do so, we shall want to have good evidence about the person's character, about possible hidden agendas, and about whether she has the sort of experiences that contribute to the kind of competence we are counting on her to have. We would also want to know that our witnesses have appropriate epistemic self-assessment,³⁰ and are not given to asserting with confidence claims that they are in a position to assert only tentatively if at all. Perhaps the most trustworthy testifiers about moral matters are those who are least inclined to offer such testimony.

Returning to Peter's case, we can see, I think, that had Peter accepted the women's testimony, his trust would have been justified. Having lived in the same household, Peter had available enough of the right kind of information to justify moving from a default position of distrust to one of trust. Both sides thought well of each other's character.

In addition, Peter had reasons to distrust his own distrust. It might be objected that Peter could not know this in advance of knowing whether the women were in fact correct. But if Peter could have known whether the women were correct, then there would have been no need to trust them in the first place. This objection forces

³⁰ For the notion of epistemic self-assessment, see Hardwig, p. 700.

us to consider whether the notion of appropriate metastance can pull enough independent weight to enter into our assessment of when trust is justified, or whether our assessment of appropriate metastance must always be consequent on our assessment of appropriate first-order trust.

I think, though, that the example itself shows why this objection is mistaken. Peter had plenty of evidence to make him suspicious of his distrust. His past track record reveals a tendency to find fewer occasions calling for antiracist and antisexist response than others. Moreover, given what Peter knew, or should have known, about how experience affects competence in these matters, he should have been aware that others might be better placed to judge than himself.

In the example, I assumed the women were in broad agreement about whether there was evidence of attitudinal sexism. Had there not been such agreement, though, it is still possible that Peter would have been justified in trusting the women who found evidence of sexism. Agreement is by no means a sure indicator of truth. Sometimes it counts against the trustworthiness of testifiers, signaling that they are in cahoots, rather than independently correct.³¹ And sometimes, especially in political contexts, it counts against the trustworthiness of testifiers by signaling that some voices have yet to be heard. In fact, given the complex interests that are so crudely captured within identity politics, unanimity among the members of an oppressed group is generally counterevidence for trustworthiness. One of the chief responsibilities of wise trusters will thus be determining whom to trust when there is disagreement.

The evidential standards that must be met before trust in someone else's moral testimony is justified are high. This means that the morally indolent can take no particular comfort from a defense of the possibility of borrowing other people's moral knowledge. There is still work to be done and responsibilities to be met before we trust testifiers on moral matters wisely. But we can see now how to respond to one of our outstanding worries, namely, the concern that accepting moral testimony conflicts with autonomy. We must subject judgments about whom we trust on these matters to rational scrutiny, and we can be called on to revise our assessments of who is trustworthy. Accepting moral testimony is shown not to be in conflict with the value of autonomy, properly understood.

We might still wonder, though, whether, had Peter relented and trusted the testimony of the women, there would have been some-

³¹ See Coady, p. 192.

thing inferior about his epistemic status. There are two questions here: Could Peter have actually come to know through trusting? Is his knowledge deficient as practical knowledge?

While I think that wise trusters can have knowledge on the basis of testimony—provided, that is, they are lucky enough to trust trustworthy knowers—I am going to leave that issue to one side so as to take up the issue of practicality.³² Even if Peter had been willing to borrow the women's moral knowledge, at the end of the day he would have been as much in the dark as to how to go about applying 'sexism' to new nonegregious cases as he ever was. Of course, we can hope that with discussion and practice he will increase his sensitivity. There is, though, no guarantee that this will happen. Borrowing may be the best that Peter can ever do.

So does this not show that there is something regrettable about the need to borrow moral knowledge? At most, it shows there are some cases in which it is regrettable. Consider now a case of a different sort, concerning ranking principles or ends. The Chinese author Luxun once quipped that the Chinese do not understand freedom, they only understand cooking oil. Let us start from this quip and construct a case. It is China in the early twentieth century and two long-time friends are discussing whether to join a political move-

³² Briefly, the argument for this claim is as follows: in accepting your testimony, I am attempting to borrow your knowledge, that is, to claim the knowledge you have as my own on the basis of your having shared it with me. But I cannot borrow from you what you do not have. Thus, before my accepting your testimony can be knowledge conferring, you yourself must have knowledge regarding the question at issue. If you are wrong, and I take your word for it, I obviously do not know it; likewise, if you are right, but only accidentally right, my being right is no more securely tethered to the truth of the matter than your being right is.

But while your knowing that p is a necessary condition for my knowing that p on the basis of your say-so, it is not a sufficient condition. Even if you know that p , you might not be trustworthy with respect to this exchange of information. For example, you might be telling me that p because you think that is a good way to have me come to believe that not- p . This would be true, if, for example, you falsely believed that I was acutely suspicious of you and so would take your word to be counterevidence for the truth of p . To the extent that my acceptance of testimony rests on trust (and not, for example, on having access to the roundabout reason you happened to tell me the truth), you must be *trustworthy* in the relevant domain before I can know that p on the basis of your telling me that p . If you are not trustworthy, then you could just as well be telling me a falsehood, and I would be in no position to discern this; thus my connection to the truth of p , arrived at through your say-so, would be too tenuous to count as giving me knowledge that p , even if you yourself knew that p .

None of this shows that those who trust wisely, and are fortunate enough to trust trustworthy knowers, can thereby have knowledge. But, as Coady has persuasively argued, our knowledge is so testimony saturated that, if the wise truster cannot know, then there is little indeed that anyone can know.

ment for democratic freedom. One has returned from study in Europe and has himself joined and is trying to persuade the other to do so. His friend is reluctant: "It is dangerous," he says. "Don't we have responsibilities to maintain certain traditions? And isn't the very idea of such change altogether too European?" The one who has returned from Europe tries to convince his friend that he does not know what freedom is really like and that, if he did, he, too, would consider the risk worth it and the tradition stifling. "You haven't," he says, "tasted freedom, and if you had, you would choose as I do." They talk, we might suppose, into the small hours of many nights. Finally, the other also decides to join the political movement, because of his trust in his friend. But that choice does not leave everything as it was before. Having made the choice, he is thrown into different circumstances and comes to see things differently. Through pursuing democratic freedom, he comes to understand more fully its value. He now knows first-hand what he initially knew only second-hand, and his belief, originally grounded in testimony, comes to have rich alternative sources of support.

It might be objected that the newcomer to the political movement did not need to trust his friend at all and that he was not in fact deferring to his friend's moral expertise—he was just trying it out and there is no deference involved in trying something out. But this reconstruction of the example overlooks two important features of it. First, the newcomer initially had objections to joining the movement, objections that had to be overcome before he could see himself as justified in even merely "trying" it out. Second, without at least some trust in the friend, it would not have been responsible for the newcomer to put himself in circumstances which would predictably lead to changes in his experiences and values, changes which from his current perspective look morally suspicious. It is no comfort to know that once we make a choice, we shall come to have experiences that confirm that choice as the right choice, unless we also believe that such confirmation will be truth conducive. Joining a cult might lead us to have experiences that will confirm our decision—but that we shall have such experiences is no reason to join; rather, it is reason to steer well clear, least our judgment come to be distorted by our membership. The newcomer thus had to trust the friend that membership in the political group would generate experiences that were truth conducive rather than error confirming.

At this point, it might be thought that my defense of testimony remains grudging. It is better when knowledge initially gained through testimony is fully integrated into a new evaluative outlook and the

person comes to know first-hand what she originally knew only second-hand. Someone might accept this and yet still maintain that we have an obligation to cultivate all our moral capacities and so use testimony concerning moral matters only as a last resort when there is no other avenue open to us to secure the knowledge we need. Is it not part of living a morally serious life that we keep pushing at the boundaries of that which we must take on trust until the place for testimony shrinks to near vanishing?

The suggestion that we have an obligation to keep pushing back the place for testimony rests on some empirical assumptions, however, that we should hold doubtful: in particular, it assumes that we can develop our moral capacities to an equal extent. Given the world of value is complex and the capacities needed to navigate in it various, there is reason to suppose that we cannot develop all morally relevant sensitivities to an equal extent. And we may do worse by trying than by practicing wise trusting.

These remarks suggest an interesting answer to the question of the unity of the virtues. I appear to be assuming that you can have one virtue without having another. Owen Flanagan³³ argues that psychological research shows the capacities that support the particular virtues do not come in a unified package; thus one can be courageous without being kind, and we should abandon the thesis of the unity of virtue. If one feels the pull of these empirical considerations, yet also feels the pull of a Socratic a priori argument for the unity of the virtues, then hope for reconciliation lies in the notion of wise trusting. Recall the argument for the unity of the virtues goes as follows: virtues are states that issue in only right action. Justice and kindness can *appear* to conflict: if one's kindness is not to issue in wrong action, one must also possess the virtue of justice (kindness that leads to injustice is no real kindness). If we think of a virtue as a sensitivity to considerations of a certain kind, then given that there are no limits on the combinations of such considerations that can be present in a choice situation, such apparent conflict between the virtues is always a possibility.³⁴ Thus, if the virtues are to result in only right actions, one cannot have one without having them all.

The notion of wise trusting lets us accept the empirical claims typically taken to support the disunity of the virtues, while retaining a kind of unity to the virtues. While the virtuous person might not her-

³³ *Varieties of Moral Personality: Ethics and Psychological Realism* (Cambridge: Harvard, 1991).

³⁴ For this formulation of the Socratic argument, see John McDowell, "Virtue and Reason," *The Monist*, LXII (1979): 331-50, especially pp. 331-33.

self have all the capacities necessary for her virtue to lead to only right action, she will have appropriate self-assessment and appreciate the limits of her capacities. She will know when she needs the moral help of others. Moreover, she will be wise in whom to trust. Her capacities, supplemented by the capacity for wise trust, will issue in only right action.

If these reflections are along the right lines, then we can see another sense in which the virtuous person's virtue is essentially social: it is not just that the social sphere provides much of the domain for the exercise of virtue, nor just that it provides the conditions for sustaining it; rather, the knowledge that the virtuous person has is knowledge made available to her in part through epistemic dependence on others.

KAREN JONES

Cornell University